

LAB 02: Advanced Feature Engineering in BQML

Learning Objectives

- Perform a feature cross using BigQuery's ML.FEATURE_CROSS
- Derive a coordinate feature
- Feature cross coordinate features
- Code cleanup
- Apply the BUCKETIZE function
- Apply the TRANSFORM clause
- Apply L2 Regularization
- Evaluate model performance

Introduction

In this lab, we continue to utilize feature engineering to improve the prediction of the fare amount for a taxi ride in New York City by reducing the RMSE.

In this Notebook, we perform a feature cross using BigQuery's ML.FEATURE_CROSS, derive coordinate features, feature cross coordinate features, cleanup the code, apply the BUCKETIZE function, the TRANSFORM clause, L2 Regularization, and evaluate model performance throughout the process.

Each learning objective will correspond to a **#TODO** in this student lab notebook -- try to complete this notebook first and then review the [solution notebook \(../solution/feateng-solution_bqml.ipynb\)](#). **NOTE TO SELF:** UPDATE HYPERLINK.

Model 4: Apply the ML.FEATURE_CROSS clause to categorical features

BigQuery ML now has ML.FEATURE_CROSS, a pre-processing clause that performs a feature cross.

- ML.FEATURE_CROSS generates a [STRUCT](https://cloud.google.com/bigquery/docs/reference/standard-sql/data-types#struct-type) (<https://cloud.google.com/bigquery/docs/reference/standard-sql/data-types#struct-type>) feature with all combinations of crossed categorical features, except for 1-degree items (the original features) and self-crossing items.
- Syntax: ML.FEATURE_CROSS(STRUCT(features), degree)
- The feature parameter is a categorical features separated by comma to be crossed. The maximum number of input features is 10. An unnamed feature is not allowed in features. Duplicates are not allowed in features.
- Degree(optional): The highest degree of all combinations. Degree should be in the range of [1, 4]. Default to 2.

Output: The function outputs a STRUCT of all combinations except for 1-degree items (the original features) and self-crossing items, with field names as concatenation of original feature names and values as the concatenation of the column string values.

INSTRUCTOR NOTE ONLY:

Why ML.Feature_Cross uses STRUCT - All elements in an array have the same size because all elements are of the same datatype whereas STRUCT can contain elements of dissimilar datatype. Hence, all elements are of different size. For example, hourofday and dayofweek are dissimilar datatypes.

Exercise: The ML.Feature_Cross statement contains errors. Correct the errors and run the query.

In []:

```
%%bigquery
#Objective: Perform a feature cross using BigQuery's ML.FEATURE_CROSS
CREATE OR REPLACE MODEL feat_eng.model_4
OPTIONS
  (model_type='linear_reg',
   input_label_cols=['fare_amount'])
AS
SELECT
  fare_amount,
  passengers,
  #pickup_datetime,
  #EXTRACT(DAYOFWEEK FROM pickup_datetime) AS dayofweek,
  #EXTRACT(HOUR FROM pickup_datetime) AS hourofday,
  #CONCAT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING),
          #CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING)) AS hourofday,
  ML.FEATURE_CROSS(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING),
                  CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  pickuplon,
  pickuplat,
  dropofflon,
  dropofflat

FROM `feat_eng.feateng_training_data`
```

In [18]:

```
%%bigquery
#SOLUTION
CREATE OR REPLACE MODEL feat_eng.model_4
OPTIONS
  (model_type='linear_reg',
   input_label_cols=['fare_amount'])
AS
SELECT
  fare_amount,
  passengers,
  #pickup_datetime,
  #EXTRACT(DAYOFWEEK FROM pickup_datetime) AS dayofweek,
  #EXTRACT(HOUR FROM pickup_datetime) AS hourofday,
  #CONCAT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING),
          #CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING)) AS hourofday,
  ML.FEATURE_CROSS(STRUCT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING) AS dayofw
eek,
  CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  pickuplon,
  pickuplat,
  dropofflon,
  dropofflat

FROM `feat_eng.feateng_training_data`
```

Out[18]:

—

Exercise: Create two SQL statements to evaluate the model.

In [19]:

```
%%bigquery
SELECT * FROM ML.EVALUATE(MODEL feat_eng.model_4)
```

Out[19]:

	mean_absolute_error	mean_squared_error	mean_squared_log_error	median_absolute_error	r2
0	5.500029	93.260577	0.27445	3.807548	0.



In [20]:

```
%%bigquery
SELECT SQRT(mean_squared_error) AS rmse FROM ML.EVALUATE(MODEL feat_eng.model_4)
```

Out[20]:

	rmse
0	9.657152

Sliding down the slope toward a loss minimum (reduced taxi fare)!

- Our fourth model above gives us an RMSE of 9.65 for estimating fares. Recall our heuristic benchmark was 8.29. This may be the result of feature crossing. Let's apply more feature engineering techniques to see if we can't get this loss metric lower!

Model 5: Feature cross coordinate features to create a Euclidean feature

Pickup coordinate:

- pickup_longitude AS pickuplon
- pickup_latitude AS pickuplat

Dropoff coordinate:

- #dropoff_longitude AS dropofflon
- #dropoff_latitude AS dropofflat

Coordinate Features:

- The pick-up and drop-off longitude and latitude data are crucial to predicting the fare amount as fare amounts in NYC taxis are largely determined by the distance traveled. As such, we need to teach the model the Euclidean distance between the pick-up and drop-off points.
- Recall that latitude and longitude allows us to specify any location on Earth using a set of coordinates. In our training data set, we restricted our data points to only pickups and drop offs within NYC. New York city has an approximate longitude range of -74.05 to -73.75 and a latitude range of 40.63 to 40.85.
- The dataset contains information regarding the pickup and drop off coordinates. However, there is no information regarding the distance between the pickup and drop off points. Therefore, we create a new feature that calculates the distance between each pair of pickup and drop off points. We can do this using the Euclidean Distance, which is the straight-line distance between any two coordinate points.
- We need to convert those coordinates into a single column of a spatial data type. We will use the ST_DISTANCE and the ST_GEOGPOINT functions.
- ST_DISTANCE: ST_DISTANCE(geography_1, geography_2). Returns the shortest distance in meters between two non-empty GEOGRAPHYs (e.g. between two spatial objects).
- ST_GEOGPOINT: ST_GEOGPOINT(longitude, latitude). Creates a GEOGRAPHY with a single point. ST_GEOGPOINT creates a point from the specified FLOAT64 longitude and latitude parameters and returns that point in a GEOGRAPHY value.

Exercise: Derive a coordinate feature.

- Convert the feature coordinates into a single column of a spatial data type. Use the The ST_Distance and the ST.GeogPoint functions. SAMPLE CODE: ST_Distance(ST_GeogPoint(value1,value2), ST_GeogPoint(value3, value4)) AS euclidean

In []:

```
# TODO: Your code goes here
# Objective: Derive a coordinate feature
```

In [21]:

```
%%bigquery
#SOLUTION
CREATE OR REPLACE MODEL feat_eng.model_5
OPTIONS
  (model_type='linear_reg',
   input_label_cols=['fare_amount'])
AS
SELECT
  fare_amount,
  passengers,
  #pickup_datetime,
  #EXTRACT(DAYOFWEEK FROM pickup_datetime) AS dayofweek,
  #EXTRACT(HOUR FROM pickup_datetime) AS hourofday,
  #CONCAT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING),
          #CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING)) AS hourofday,
  ML.FEATURE_CROSS(STRUCT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING) AS dayof
week,
  CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  #pickuplon,
  #pickuplat,
  #dropofflon,
  #dropofflat,
  ST_Distance(ST_GeogPoint(pickuplon, pickuplat), ST_GeogPoint(dropofflon, dropofflat))
  AS euclidean

FROM `feat_eng.feateeng_training_data`
```

Out[21]:

—

Exercise: Create two SQL statements to evaluate the model.

In []:

```
# TODO: Your code goes here
```

In [22]:

```
%%bigquery
SELECT * FROM ML.EVALUATE(MODEL feat_eng.model_5)
```

Out[22]:

	mean_absolute_error	mean_squared_error	mean_squared_log_error	median_absolute_error	r2
0	3.121311	31.229717	0.106923	2.218179	0.



In [23]:

```
%%bigquery
SELECT SQRT(mean_squared_error) AS rmse FROM ML.EVALUATE(MODEL feat_eng.model_5)
```

Out[23]:

	rmse
0	5.588356

Model 6: Feature cross pick-up and drop-off locations features

In this section, we feature cross the pick-up and drop-off locations so that the model can learn pick-up-drop-off pairs that will require tolls.

This step takes the geographic point corresponding to the pickup point and grids to a 0.1-degree-latitude/longitude grid (approximately 8km x 11km in New York—we should experiment with finer resolution grids as well). Then, it concatenates the pickup and dropoff grid points to learn “corrections” beyond the Euclidean distance associated with pairs of pickup and dropoff locations.

Because the lat and lon by themselves don't have meaning, but only in conjunction, it may be useful to treat the fields as a pair instead of just using them as numeric values. However, lat and lon are continuous numbers, so we have to discretize them first. That's what SnapToGrid does.

- ST_SNAPTOGRID: ST_SNAPTOGRID(geography_expression, grid_size). Returns the input GEOGRAPHY, where each vertex has been snapped to a longitude/latitude grid. The grid size is determined by the grid_size parameter which is given in degrees.

REMINDER: The ST_GEOGPOINT creates a GEOGRAPHY with a single point. ST_GEOGPOINT creates a point from the specified FLOAT64 longitude and latitude parameters and returns that point in a GEOGRAPHY value. The ST_Distance function returns the minimum distance between two spatial objects. It also returns meters for geographies and SRID units for geometries.

Excercise: The following SQL statement is incorrect. Modify the code to feature cross the pick-up and drop-off locations features.

In []:

```

%%bigquery
#Objective: Feature cross coordinate features
CREATE OR REPLACE MODEL feat_eng.model_6
OPTIONS
  (model_type='linear_reg',
   input_label_cols=['fare_amount'])
AS
SELECT
  fare_amount,
  passengers,
  #pickup_datetime,
  #EXTRACT(DAYOFWEEK FROM pickup_datetime) AS dayofweek,
  #EXTRACT(HOUR FROM pickup_datetime) AS hourofday,
  #CONCAT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING),
          #CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING)) AS hourofday,
  ML.FEATURE_CROSS(STRUCT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING) AS dayof
week,
  CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  #pickuplon,
  #pickuplat,
  #dropofflon,
  #dropofflat,
  ST_AsText(ST_SnapToGrid(ST_GeogPoint(pickuplat,pickuplon,pickuplat), 0.05)),
  ST_AsText(ST_GeogPoint(dropofflon, dropofflat,dropofflon), 0.04)) AS pickup_and_dropoff

FROM `feat_eng.feateng_training_data`

```


In [142]:

```
%%bigquery
#SOLUTION
CREATE OR REPLACE MODEL feat_eng.model_6
OPTIONS
  (model_type='linear_reg',
   input_label_cols=['fare_amount'])
AS
SELECT
  fare_amount,
  passengers,
  #pickup_datetime,
  #EXTRACT(DAYOFWEEK FROM pickup_datetime) AS dayofweek,
  #EXTRACT(HOUR FROM pickup_datetime) AS hourofday,
  ML.FEATURE_CROSS(STRUCT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING) AS dayof
week,
  CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  #pickuplon,
  #pickuplat,
  #dropofflon,
  #dropofflat,
  ST_Distance(ST_GeogPoint(pickuplon, pickuplat), ST_GeogPoint(dropofflon, dropofflat))
  AS euclidean,
  CONCAT(ST_AsText(ST_SnapToGrid(ST_GeogPoint(pickuplon, pickuplat), 0.01)),
  ST_AsText(ST_SnapToGrid(ST_GeogPoint(dropofflon, dropofflat), 0.01)))
  AS pickup_and_dropoff

FROM `feat_eng.feateng_training_data`
```

Out[142]:

—

Exercise: Create two SQL statements to evaluate the model.

In [7]:

```
%%bigquery
#SOLUTION
SELECT * FROM ML.EVALUATE(MODEL feat_eng.model_6)
```

Out[7]:

	mean_absolute_error	mean_squared_error	mean_squared_log_error	median_absolute_error	r2
0	2.678752	34.885692	0.088598	1.485065	0.



In [144]:

```
%%bigquery
SELECT SQRT(mean_squared_error) AS rmse FROM ML.EVALUATE(MODEL feat_eng.model_6)
```

Out[144]:

	rmse
0	5.906411

Code Clean Up

Exercise: Clean up the code to see where we are

Remove all the commented statements in the SQL statement. We should now have a total of five input features for our model.

1. fare_amount
2. passengers
3. day_hr
4. euclidean
5. pickup_and_dropoff

In []:

```
%%bigquery
#Objective: Code Clean Up
CREATE OR REPLACE MODEL feat_eng.model_6
OPTIONS
  (model_type='linear_reg',
   input_label_cols=['fare_amount'])
AS
SELECT
  fare_amount,
  passengers,
  #pickup_datetime,
  #EXTRACT(DAYOFWEEK FROM pickup_datetime) AS dayofweek,
  #EXTRACT(HOUR FROM pickup_datetime) AS hourofday,
  ML.FEATURE_CROSS(STRUCT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING) AS dayof
week,
  CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  #pickuplon,
  #pickuplat,
  #dropofflon,
  #dropofflat,
  ST_Distance(ST_GeogPoint(pickuplon, pickuplat), ST_GeogPoint(dropofflon, dropofflat))
  AS euclidean,
  CONCAT(ST_AsText(ST_SnapToGrid(ST_GeogPoint(pickuplon, pickuplat), 0.01)),
    ST_AsText(ST_SnapToGrid(ST_GeogPoint(dropofflon, dropofflat), 0.01)))
  AS pickup_and_dropoff

FROM `feat_eng.feateng_training_data`
```

In [24]:

```
%%bigquery
#SOLUTION
CREATE OR REPLACE MODEL feat_eng.model_6
OPTIONS
  (model_type='linear_reg',
   input_label_cols=['fare_amount'])
AS
SELECT
  fare_amount,
  passengers,
  ML.FEATURE_CROSS(STRUCT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING) AS dayof
week,
  CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  ST_Distance(ST_GeogPoint(pickuplon, pickuplat), ST_GeogPoint(dropofflon, dropofflat))
  AS euclidean,
  CONCAT(ST_AsText(ST_SnapToGrid(ST_GeogPoint(pickuplon, pickuplat), 0.01)),
    ST_AsText(ST_SnapToGrid(ST_GeogPoint(dropofflon, dropofflat), 0.01)))
  AS pickup_and_dropoff

FROM `feat_eng.feateng_training_data`
```

Out[24]:

—

BQML's Pre-processing functions:

Here are some of the preprocessing functions in BigQuery ML:

- `ML.FEATURE_CROSS(STRUCT(features))` does a feature cross of all the combinations
- `ML.POLYNOMIAL_EXPAND(STRUCT(features), degree)` creates x , x^2 , x^3 , etc.
- `ML.BUCKETIZE(f, split_points)` where `split_points` is an array

Model 7: Apply the BUCKETIZE Function

BUCKETIZE

Bucketize is a pre-processing function that creates "buckets" (e.g bins) - e.g. it bucketizes a continuous numerical feature into a string feature with bucket names as the value.

- `ML.BUCKETIZE(feature, split_points)`
- `feature`: A numerical column.
- `split_points`: Array of numerical points to split the continuous values in feature into buckets. With n split points ($s_1, s_2 \dots s_n$), there will be $n+1$ buckets generated.
- `Output`: The function outputs a `STRING` for each row, which is the bucket name. *bucketname is in the format of bin*, where `bucket_number` starts from 1.
- Currently, our model uses the `ST_GeogPoint` function to derive the pickup and dropoff feature. In this lab, we use the `BUCKETIZE` function to create the pickup and dropoff feature.

Exercise: Apply the BUCKETIZE function.

- Hint: Create a `model_7`.

In []:

```
# TODO: Your code goes here
# Objective: Apply the BUCKETIZE function
```

In [1]:

```
%%bigquery
#SOLUTION
CREATE OR REPLACE MODEL feat_eng.model_7
OPTIONS
  (model_type='linear_reg',
   input_label_cols=['fare_amount'])
AS
SELECT
  fare_amount,
  passengers,
  ST_Distance(ST_GeogPoint(pickuplon, pickuplat), ST_GeogPoint(dropofflon, dropofflat)) AS
euclidean,
  ML.FEATURE_CROSS(STRUCT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING) AS day
ofweek,
  CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  CONCAT(
    ML.BUCKETIZE(pickuplon, GENERATE_ARRAY(-78, -70, 0.01)),
    ML.BUCKETIZE(pickuplat, GENERATE_ARRAY(37, 45, 0.01)),
    ML.BUCKETIZE(dropofflon, GENERATE_ARRAY(-78, -70, 0.01)),
    ML.BUCKETIZE(dropofflat, GENERATE_ARRAY(37, 45, 0.01))
  ) AS pickup_and_dropoff
FROM `feat_eng.feateng_training_data`
```

Out[1]:

—

Exercise: Create three SQL statements to EVALUATE the model.

In [2]:

```
%%bigquery
SELECT *, Sqrt(loss) AS rmse FROM ML.TRAINING_INFO(MODEL feat_eng.model_7)
#NOTE that with each iteration, loss decreases.
```

Out[2]:

	training_run	iteration	loss	eval_loss	learning_rate	duration_ms	rmse
0	0	3	11.231835	33.172360	0.4	6434	3.351393
1	0	2	13.263899	35.919936	0.8	5898	3.641964
2	0	1	20.108475	36.559006	0.4	6533	4.484247
3	0	0	70.490524	81.719679	0.2	3746	8.395864

In [3]:

```
%%bigquery
SELECT * FROM ML.EVALUATE(MODEL feat_eng.model_7)
```

Out[3]:

	mean_absolute_error	mean_squared_error	mean_squared_log_error	median_absolute_error	r2
0	2.674187	33.17236	0.089913	1.51821	0.



In [4]:

```
%%bigquery
SELECT SQRT(mean_squared_error) AS rmse FROM ML.EVALUATE(MODEL feat_eng.model_7)
```

Out[4]:

	rmse
0	5.759545

Final Model: Apply the TRANSFORM clause and L2 Regularization

Before we perform our prediction, we should encapsulate the entire feature set in a TRANSFORM clause. BigQuery ML now supports defining data transformations during model creation, which will be automatically applied during prediction and evaluation. This is done through the TRANSFORM clause in the existing CREATE MODEL statement. By using the TRANSFORM clause, user specified transforms during training will be automatically applied during model serving (prediction, evaluation, etc.)

In our case, we are using the TRANSFORM clause to separate out the raw input data from the TRANSFORMED features. The input columns of the TRANSFORM clause is the query_expr (AS SELECT part). The output columns of TRANSFORM from select_list are used in training. These transformed columns are post-processed with standardization for numerics and one-hot encoding for categorical variables by default.

The advantage of encapsulating features in the TRANSFORM clause is the client code doing the PREDICT doesn't change, e.g. our model improvement is transparent to client code. Note that the TRANSFORM clause MUST be placed after the CREATE statement.

[L2 Regularization \(https://developers.google.com/machine-learning/glossary/#L2_regularization\)](https://developers.google.com/machine-learning/glossary/#L2_regularization)

Sometimes, the training RMSE is quite reasonable, but the evaluation RMSE illustrate more error. Given the severity of the delta between the EVALUATION RMSE and the TRAINING RMSE, it may be an indication of overfitting. When we do feature crosses, we run into the risk of overfitting (for example, when a particular day-hour combo doesn't have enough taxi rides).

Overfitting is a phenomenon that occurs when a machine learning or statistics model is tailored to a particular dataset and is unable to generalize to other datasets. This usually happens in complex models, like deep neural networks. Regularization is a process of introducing additional information in order to prevent overfitting.

Therefore, we will apply L2 Regularization to the final model. As a reminder, a regression model that uses the L1 regularization technique is called Lasso Regression while a regression model that uses the L2 Regularization technique is called Ridge Regression. The key difference between these two is the penalty term. Lasso shrinks the less important feature's coefficient to zero, thus removing some features altogether. Ridge regression adds "squared magnitude" of coefficient as a penalty term to the loss function.

In other words, L1 limits the size of the coefficients. L1 can yield sparse models (i.e. models with few coefficients); Some coefficients can become zero and eliminated.

L2 regularization adds an L2 penalty equal to the square of the magnitude of coefficients. L2 will not yield sparse models and all coefficients are shrunk by the same factor (none are eliminated).

The regularization terms are 'constraints' by which an optimization algorithm must 'adhere to' when minimizing the loss function, apart from having to minimize the error between the true y and the predicted \hat{y} . This in turn reduces model complexity, making our model simpler. A simpler model can reduce the chances of overfitting.

Exercise: Apply the TRANSFORM clause and L2 Regularization to the final model and run the query.

In []:

```
# TODO: Your code goes here
# Objective: Apply the TRANSFORM clause and L2 Regularization to the final model and run the query.
```

In [12]:

```
%%bigquery
#SOLUTION
CREATE OR REPLACE MODEL feat_eng.final_model
TRANSFORM(
  fare_amount,
  #SQRT( (pickuplon-dropofflon)*(pickuplon-dropofflon) + (pickuplat-dropofflat)*(pickuplat-dropofflat) ) AS euclidean,
  ST_Distance(ST_GeogPoint(pickuplon, pickuplat), ST_GeogPoint(dropofflon, dropofflat)) AS euclidean,
  ML.FEATURE_CROSS(STRUCT(CAST(EXTRACT(DAYOFWEEK FROM pickup_datetime) AS STRING) AS dayofweek,
    CAST(EXTRACT(HOUR FROM pickup_datetime) AS STRING) AS hourofday)) AS day_hr,
  CONCAT(
    ML.BUCKETIZE(pickuplon, GENERATE_ARRAY(-78, -70, 0.01)),
    ML.BUCKETIZE(pickuplat, GENERATE_ARRAY(37, 45, 0.01)),
    ML.BUCKETIZE(dropofflon, GENERATE_ARRAY(-78, -70, 0.01)),
    ML.BUCKETIZE(dropofflat, GENERATE_ARRAY(37, 45, 0.01))
  ) AS pickup_and_dropoff
)
OPTIONS(input_label_cols=['fare_amount'], model_type='linear_reg', l2_reg=0.1)
AS

SELECT * FROM feat_eng.feateng_training_data
```

Out[12]:

—

Exercise: Create three SQL statements to EVALUATE the final model.

In []:

```
#HINT: You will need to create three separate cells for your SQL queries.
```

In [7]:

```
%%bigquery
SELECT *, SQRT(loss) AS rmse FROM ML.TRAINING_INFO(MODEL feat_eng.final_model)
```

Out[7]:

training_run	iteration	loss	eval_loss	learning_rate	duration_ms	rmse	
0	0	0	11.120628	21.654845	None	232685	3.334761

In [8]:

```
%%bigquery
SELECT * FROM ML.EVALUATE(MODEL feat_eng.final_model)
```

Out[8]:

	mean_absolute_error	mean_squared_error	mean_squared_log_error	median_absolute_error	r2
0	2.262816	21.654845	0.068715	1.358259	0.



In [9]:

```
%%bigquery
SELECT SQRT(mean_squared_error) AS rmse FROM ML.EVALUATE(MODEL feat_eng.final_model)
```

Out[9]:

	rmse
0	4.653477